POLITECNICO

MILANO 1863

Clustering PM$_{10}$ and other cute stuff

PROJECT REPORT OF BAYESIAN STATISTICS - MATHEMATICAL ENGINEERING

Giulia Mezzadri, Ettore Modina, Oswaldo Jesus Morales Lopez, Federico Angelo Mor, Abylaikhan Orynbassar, Federica Rena

Professor:
Alessandra Guglielmi

Tutors:
Michela Frigeri
Alessandro Carminati

Academic year:
2023-2024

Abstract[a] DONE In this project we undertake a comprehensive clustering analysis of PM10 levels in the Lombardy region (Italy), employing four different Bayesian models to account for the complex nature of our data, which comprise spatio-temporal measurements of PM10, together with many other environmental variables, collected from various monitoring stations displaced across the entire region. The main objective was to leverage on covariates, station locations and time trends to cluster weekly PM10 data over a one-year period. Our analysis revealed distinct clusters for each time step, with a noteworthy influence of morphological terrain characteristics (e.g. altitude, wind speed) and anthropological factors (e.g. agricultural activities, vehicles and road transports, etc.).

The analysis was executed concurrently across a set of four models, to study the different interactions and combinations of spatio-temporal aspects and covariates information. Despite some variations among the models, that however highlighted peculiar patterns and characteristics which each model independently dwelt on, a unanimous consensus emerged regarding the overall division between the stations. This study contributes valuable insights into the delicate interaction of spatial, temporal, and covariate variables in shaping PM10 levels, providing a robust foundation for understanding the clustering dynamics in the Lombardy region.

---

[a]See `https://github.com/federicomor/progetto-bayesian` for all the project codes and `https://federicomor.github.io/assets/figures/visualize.html` for the visualization page.

1. Introduction

[Com17] Particulate matter with a diameter of 10 micrometers or less, known as PM10, comprises small airborne particles sourced from various origins, posing potential health risks upon inhalation due to their ability to deeply penetrate the respiratory system. The meticulous monitoring of PM10 levels is imperative for comprehensive air quality assessment and the safeguarding of public health.

This paper embarks on a project with the overarching aim of identifying both natural and anthropogenic factors contributing to elevated PM10 levels. Employing a clustering analysis, our objective is to delineate distinct regions within Lombardy, unveiling discernible patterns influencing particulate levels.

Drawing upon data from the Agrimonia project, which encompasses diverse measurements, our focus centers on weekly averages across a one-year timeframe. Our analytical approach involves the utilization of various models, including DRPM and SPPM, alongside additional models for covariate selection.

In subsequent sections, we delve into the dataset cleaning process, present individual analyses for each model, and expound on our interpretation of results. Visualization plays a pivotal role in our exploration, with a particular emphasis on manual (visual) interpretation to extract nuanced insights.

It is essential to acknowledge the inherent limitations of our approach. Notably, our lack of technical expertise in the phenomenon necessitated a wholly data-driven analysis, underscoring the importance of contextualizing our findings within this parameter.

2. The dataset

The Agrimonia dataset, developed in [FRFM$^+$23], spans from January 1, 2016, to December 31, 2021, recording observations from a network of 141 stations in the surroundings of Lombardy region. The dataset gathers measurements from five different covariate groups: air quality (AQ), weather and climate (WE), pollutants' emissions (EM), livestock (LI), land and soil characteristics (LA). In total there are 38 covariates, with our target variable lying in the AQ group, namely AQ_pm10.
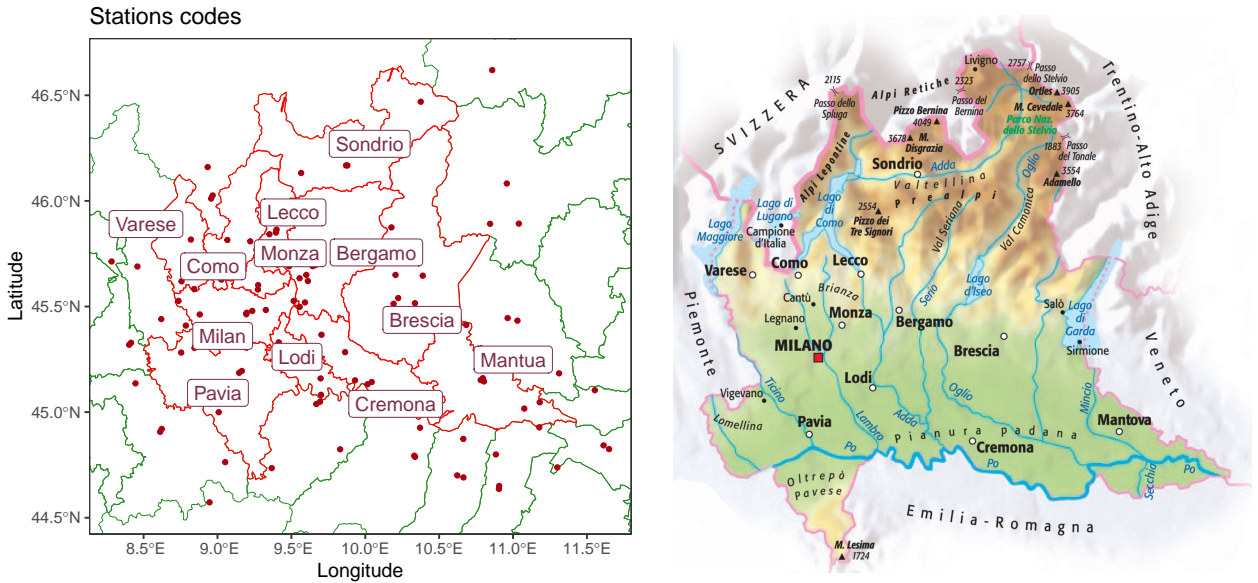

Figure 1: On the left, the map of the 105 selected stations after the data preprocessing. On the right, the geographical map on the Lombardy region, useful to understand recall the morphological structure of that area. anyway for now these are jsut provisional images, we have to think what we like to put here for example. the next ones are more definitive instead

2.1. Data inspection

2.2. Data preprocessing DONE

The goal proposed for our project was "clustering weekly data of one year of PM$_{10}$" and as such we started by selecting the year and then dividing the dataset by weeks, since originally it consisted of daily recordings. One main concern for the year selection was the presence of missing data (NA) both in the target variable and in many other ones, as we can see in Figure 2.

About the covariates, there was a considerable scarcity in the AQ group, so we were forced to remove them and save only our target PM$_{10}$ variable. This may look as a relevant information loss, since the other pollutants like PM$_{2.5}$, SO$_2$ or NH$_3$ could have related well to the PM$_{10}$ concentrations. But actually another group of variables, the EM one, stores information about those pollutants, so we still preserved the information that they carried. We also removed variable LA_soil_use, which was considerably empty for most of the stations, and after this procedure we remained with 36 variables (of which 31 were covariates) from the original 43.

Regarding the PM$_{10}$ levels, instead, there were many stations which were totally lacking of any recorded value, as revealed by Figure3; therefore we removed them and we were left with 105 stations out of the original 141, which is still a very representative set. After this cleaning procedure, a natural choice for the year would have

Figure 2: Heatmap of the missing data of all the covariates in the available dataset, i.e. considering all the six years data. On the rows there are the covariates, on the columns all the original 141 stations.

been to select the most recent one, but due to 2021 showing an increase in the missing values and being it still close to the covid-affected period after 2019, we decided in the end to choose 2018, hoping that the present years, somehow recovered from the pandemic anomalous levels, would be similar to that one.

Then we moved to the task of the weekly averaging. For the covariates, the selected year showed to be almost full of values except for three stations, which were lacking of all values in three variables of the LI and LA groups. Since this was a small problematic case and concerned stations outside the Lombardy area, the one of primary interest, we didn't deem necessary to remove them completely, but instead we filled in the missing data with an average of the values of the three closest stations on the map.

Also for the PM_{10} data there were still some missing spots, but now affecting almost all stations. Initially we thought of filling them by using, for each station, a linear interpolation between the closest-in-time present data


Figure 3: Heatmap of the missing data (in white) of the PM_{10} values recorded in the available dataset. On the rows there are all the original 141 stations, on the columns all the 2192 days composing the six years.

around a set of missing ones. This would have allowed the build of the weekly division by simply averaging over those (now all complete) values. But we thought that this method would have induced a double approximation: the first one in the NA filling and the second one in the weekly averaging. So in the end we decided to directly build the weekly division by averaging not necessary on the complete set of seven days, but just on the available values in a given week. We did this procedure on all the numeric variables, as well as on the categorical ones (e.g. the wind direction) but using the mode instead of the mean.

In this way we got the final dataset, on which we then performed a logarithmic transformation to the PM_{10} variable, to achieve a normal distribution, followed by a shift to bring them into having zero mean. We also standardized the numerical covariates, including the spatial coordinates. This allowed us to enhance the suitability of the data for the subsequent statistical models, which for example assumed a normal distribution of the target data, and in general worked better using centered data, to accommodate the prior distribution support of the parameters. This comprehensive processing dataset formed the foundation for our investigation into the factors influencing PM10 levels in the Lombardy region.

3.   Models DONE

For our analysis we looked into models which could tailor the complex nature of our data, exploiting spatial and temporal information, together with covariates, with a clustering target in mind. Unfortunately, there was no "holy grail" which could manage to harness all those levels of information, but nonetheless we found four models which in the end worked well for our task.

We will now see them in details, but for a clear preview of their characteristics refer to Table 1.

model name	Time	Space	Covariates
sPPM	✗	✓	✗
DRPM	✓	✓	✗
Gaussian PPMx	✗	✗	✓
Curve PPMx	✗	✗	✓

Table 1: Summary of the functional characteristics of the models at hand. All of them were able to perform clustering natively.

3.1.   sPPM model

3.2.   DRPM model DONE

The second model we focused on, outside of the `PPMSuite` package, is the Dependent Modeling of Temporal Sequences of Random Partitions (DRPM), developed in [PQD22]. The main objective of the authors was to define a spatio-temporal model capable of performing "smooth" clusterings, i.e. a framework which would favour a gentle evolution in time of the units allocations, rather than abrupt (and therefore less interpretable) changes in them. This result was clearly reached also in our analysis, as we will describe more precisely in section 4.2, where we witnessed a more regular trend in the clusters definition for the DRPM model with respect to the other ones.

The model that we used, fully detailed in (1), starts by assuming a first order dependence relation between clusters, meaning that the conditional distribution of ρ_t given $\rho_{t-1}, \ldots, \rho_1$ just depends on ρ_{t-1}. This idea is implemented using a temporal dependence parameter $\alpha \in [0, 1]$ which controls the level of flexibility in the cluster allocation variables: the higher is α, the higher is the tendency of units to remain in their current cluster, meaning that clusters ρ_{t+1} will be similar to ρ_t. Conversely, when α approaches 0, we would get more independent clusters. In this way the clusters allocations variables c_t will follow a temporal Random Partition Model (the entry tRPM in the model formulation) driven by the sequence of α_t and the Dirichlet dispersion parameter M.

$$Y_{it}|Y_{it-1}, \boldsymbol{\mu}_t^\star, \boldsymbol{\sigma}_t^{2\star}, \boldsymbol{\eta}, \boldsymbol{c}_t \overset{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{it}t}^\star + \eta_{1i}Y_{it-1}, \sigma_{c_{it}t}^{2\star}(1-\eta_{1i}^2)) \quad i = 1, \dots, n \quad \text{and} \quad t = 2, \dots, T$$

$$Y_{i1} \overset{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{i1}1}^\star, \sigma_{c_{i1}1}^{2\star})$$

$$\xi_i = \text{Logit}(\tfrac{1}{2}(\eta_{1i}+1)) \overset{\text{ind}}{\sim} \text{Laplace}(a,b)$$

$$(\mu_{jt}^\star, \sigma_{jt}^\star) \overset{\text{ind}}{\sim} \mathcal{N}(\theta_t, \tau_t^2) \times \mathcal{U}(0, A_\sigma) \tag{1}$$

$$\theta_t|\theta_{t-1} \overset{\text{ind}}{\sim} \mathcal{N}((1-\phi_1)\phi_0 + \phi_1\theta_{t-1}, \lambda^2(1-\phi_1^2))$$

$$(\theta_1, \tau_t) \sim \mathcal{N}(\phi_0, \lambda^2) \times \mathcal{U}(0, A_\tau)$$

$$(\phi_0, \phi_1, \lambda) \sim \mathcal{N}(m_0, s_0^2) \times \mathcal{U}(-1, 1) \times \mathcal{U}(0, A_\lambda)$$

$$\{\boldsymbol{c}_t, \dots, \boldsymbol{c}_T\} \sim \text{tRPM}(\boldsymbol{\alpha}, M) \quad \text{with} \quad \alpha_t \overset{\text{iid}}{\sim} \text{Beta}(a_\alpha, b_\alpha)$$

About the target variable Y_{it}, they modelled it with a Normal law with mean $\boldsymbol{\mu}_t^\star$ and variance $\boldsymbol{\sigma}_t^{2\star}$. The mean of that distribution actually incorporates a more sophisticated modelling introducing an autoregressive part both at the observations and at the parameters (or "atoms") level. Indeed, the Y_{it} depend on Y_{it-1} through the parameter η_{i1}, while for the $\mu_{jt}^\star$ the autoregressive structure is inside the parameter θ_t which enters in his prior distribution definition.

This deepening level allowed us to test different subsets of models and to select the best one which would suit our data. Through their package `drpm` on `R`, we fitted 8 different models based on the binary choices available for those three key parameters: the α could be set constant or varying in time, while the η_{1i} and ϕ_1 could be present (therefore introducing the autoregressive design) or not.

	model name			LPML	WAIC
model	η:No	ϕ:Yes	α_t:Yes	**1077.64**	**-2366.48**
model	η:No	ϕ:No	α_t:Yes	950.17	-2117.36
model	η:Yes	ϕ:No	α_t:No	724.34	-1474.02
model	η:No	ϕ:Yes	α_t:No	693.04	-1458.70
model	η:Yes	ϕ:No	α_t:Yes	605.32	-1287.13
model	η:No	ϕ:No	α_t:No	504.41	-1129.83
model	η:Yes	ϕ:Yes	α_t:No	445.16	-913.62
model	η:Yes	ϕ:Yes	α_t:Yes	403.05	-1264.03

Table 2: Metrics values computed for the DRPM model selection, sorted by best to worst. Higher LPML and lower WAIC values denote a better fit.

According to those tests, the best model for our scenario turned out to be the one using a time specific α and with an autoregressive component just at the atoms level, while not for the observations. Surprisingly, the model at his full complexity scored last in the ranking. We then we ran another fit on the best model, using some further refined parameters in terms of samples collection, to get the definitive results. Each fit of the 8 models tested above took around one hour, while the final fit took a little more than two hours and we ran 100000 iterations, discarding the first 60000, and thinning by 40; thus getting 1000 iterates. The high value of burn in was deemed necessary after seeing some significant oscillations even after a lot of iterations, while the thinning value was suggested by the authors and confirmed by the good trend of almost all our ACF plots (see Appendix B for them).

3.3. Gaussian PPMx model

As our third model, we decided to implement a Gaussian PPMx model, developed in `PPMSuite` package. The model can be defined as follows:

$$Y_i|\mu_j^*, \sigma_j^*, S_i \overset{\text{iid}}{\sim} N(\mu_{S_i}^*, \sigma_{S_i}^{*2}), \quad i = 1, \ldots, m$$

$$\sigma_j^* \sim U(0, A),$$

$$\mu_j^*|\mu_0, \sigma_0^2 \overset{\text{iid}}{\sim} N(\mu_0, \sigma_0^2), \quad j = 1, \ldots, k_m$$

$$\sigma_0 \sim U(0, A_0),$$

$$\mu_0 \sim N(m, s^2),$$

$$\Pr(\rho_m|x) \propto \prod_{j=1}^{K_m} c(S_j) g(x_j^*)$$

For our analysis, we decided to include covariates in the model. In order to make the results as interpretable as possible, we considered a subset of five variables that maximized the LPML. Specifically, before obtaining the final model, we considered preliminary models aimed solely at variable selection.

Firstly, we compared 31 different models, each characterized by a single variable. The first variable selected, was the one corresponding to the model with the highest LPML: *Altitude*. Secondly, we considered 30 different models, each characterized by the first chosen variable and another variable present in the dataset. The second variable chosen was the one corresponding to the model with two covariates with the highest LPML: `EM_nox_sum`. By repeating this procedure three more times, we obtained the best subset of five variables in terms of LPML, consisting of: *Altitude, EM_ nox_ sum, WE_ mode_ wind_ direction_ 100m, WE_ wind_ speed_ 100m_ max*, and *LA_ lvi* .

Once we finalized the set of covariates to use in the model, we focused on the choice of priors and other required parameters.

In particular, following comparisons among different models and observing the trace plots and autocorrelation plots, we considered the Auxilliary *N-NIG* similarity function, as suggested in the paper (to be cited) and the priors were initialized as follows:

- $m = 0$
- $s^2 = 100$
- $A = 0.5$
- $A_0 = 1$

At the end, 4300 MCMC iterates were collected, with a burn-in of 300.

In this manner, we achieved favorable trends in nearly all of our ACF plots (refer to Appendix B for them), thus reinforcing the value of our efforts, the decisions made, and the outcomes obtained.

3.4. Curve PPMx model

The last model that we considered was the Functional Gaussian PPMx model which was implemented in `PPMSuite` package using the function `curve_ppmx` function. This function applies a hierarchical functional data model, wherein B-spline coefficients undergo clustering utilizing either a PPM or a PPMx prior approach on partitions using Gaussian PPMx approach. The details of latter model is provided in section 3.3. The PPM and PPMx priors are used to group similar functional curves together based on their characteristics. The PPMx prior integrates the concept that individuals sharing similar covariate values are more likely to be grouped together (reference to paper "Predictions Based on the Clustering of..."). Since the model uses Gaussain PPMx model approach, we ran the model with the same covariates mentioned in sections 3.3.

3.5. Linear Model

We thought it could be useful to implement also a simpler baseline model to use for comparison, variable selection and to better understand the data, that could allow to try a vast range of methods faster than more complex models.

Being simpler and allowing to include covariates, we chose the linear approach, actually implementing a linear model for each station. The models considered the numerical covariates linearly but tried to allow more variability on the time considering also its sine, cosine and square.

The first idea was to try a clustering on the linear model, maybe grouping together the stations according to their betas but it was soon discarded as we deemed it redundant, since we already had different models more fitted for clustering, and to focus on variable selection instead. The model was firstly implemented through jags to try and use a selection method since with so many covariates it would have been extremely long and

computationally heavy to try methods based on partitions of covariates or spike and slab. The method returned a matrix with confidence quantiles on the columns and the covariates on the rows, with value 1 if the covariate was kept at that confidence, 0 otherwise. Our idea was to choose a quantile, keep only the corresponding column as a vector, and sum element-wise for all stations, expecting an higher value in correspondence of a useful covariate to select, and a significantly lower for covariates discarded by most stations. Unfortunately this first attempt did not lead us to any solid conclusion since, even changing the threshold and hyperparameters, the final vector had very similar values on all covariates, usually between 50 and 60, suggesting all had been selected only for about half the stations, and none was significantly more important than the others.

We then tried bayesian lasso and horseshoe methods, using the corresponding R packages and the same approach as before, since both methods returned a binary vector indicating whether or not to keep a variable. Horseshoe analysis was inconclusive, discarding all covariates, while lasso showed a great weight on the total precipitations, and a more moderate but still interesting on the livestock, lvi and hvi (related to total green area per unit horizontal ground surface area for low and high vegetation type) variables.

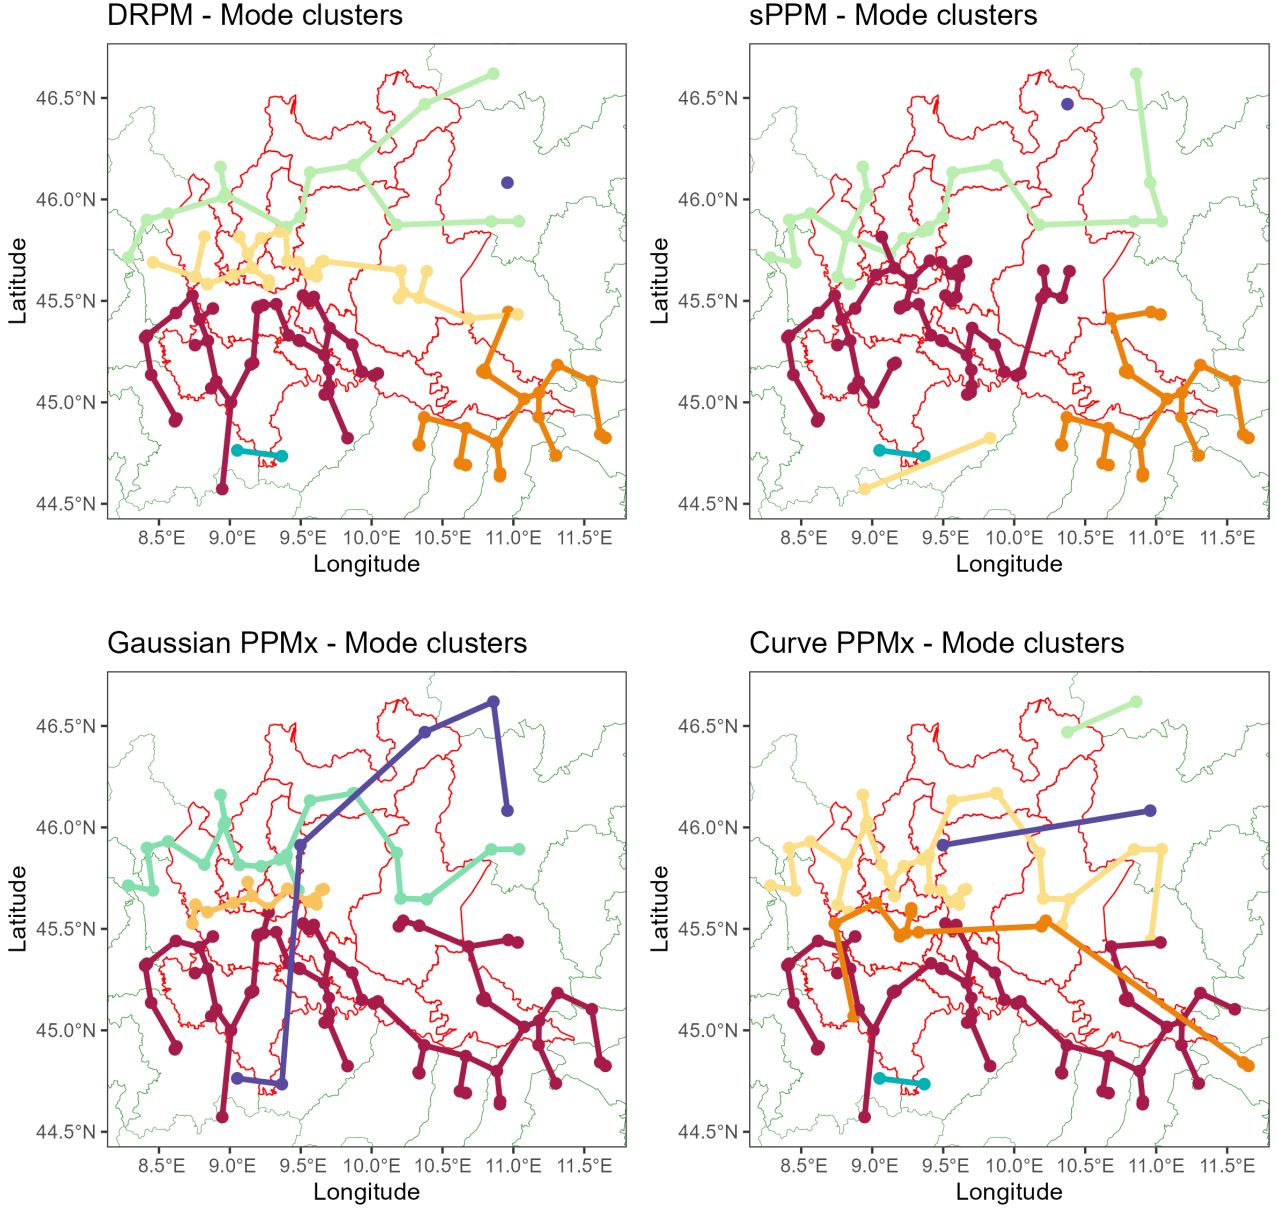# 4. Models comparison

4.1. Cluster trends

Figure 4: Maps of the most frequent clusters, throughout the 53 weeks of 2018, for all the models. See `https://federicomor.github.io/assets/figures/visualize.html` for a more detailed analysis of the plots (e.g. for all the week by week clusterings).

4.2. ARI metric ALMOST DONE

A more numerical way to compare the clustering results is through the Adjusted Random Index (ARI), developed by [HA85], which is a sort of correlation index which measures the similarity between clusterings. Given two partitions ρ_1 and ρ_2, the $\mathrm{ARI}(\rho_1, \rho_2)$ describes the amount of accordance between them, i.e. the level of agreement that they show in clustering the data. The ARI values are bounded above by one, which refers to a perfect alignment, and have zero expected value, which refers to the case of comparing two random generated partitions.

This metric was developed as a correction of the Random Index (RI) to take into account the fact the some concordance can happen by chance. While this correction is outside the scope of this analysis, the idea behind the definition of the original index lies in computing the frequency of agreements for any possible pair of units. Therefore we get $\mathrm{RI}(\rho_1, \rho_2) = (a + b)/\binom{n}{2}$, where a is the number of pairs allocated in the same subset (i.e.
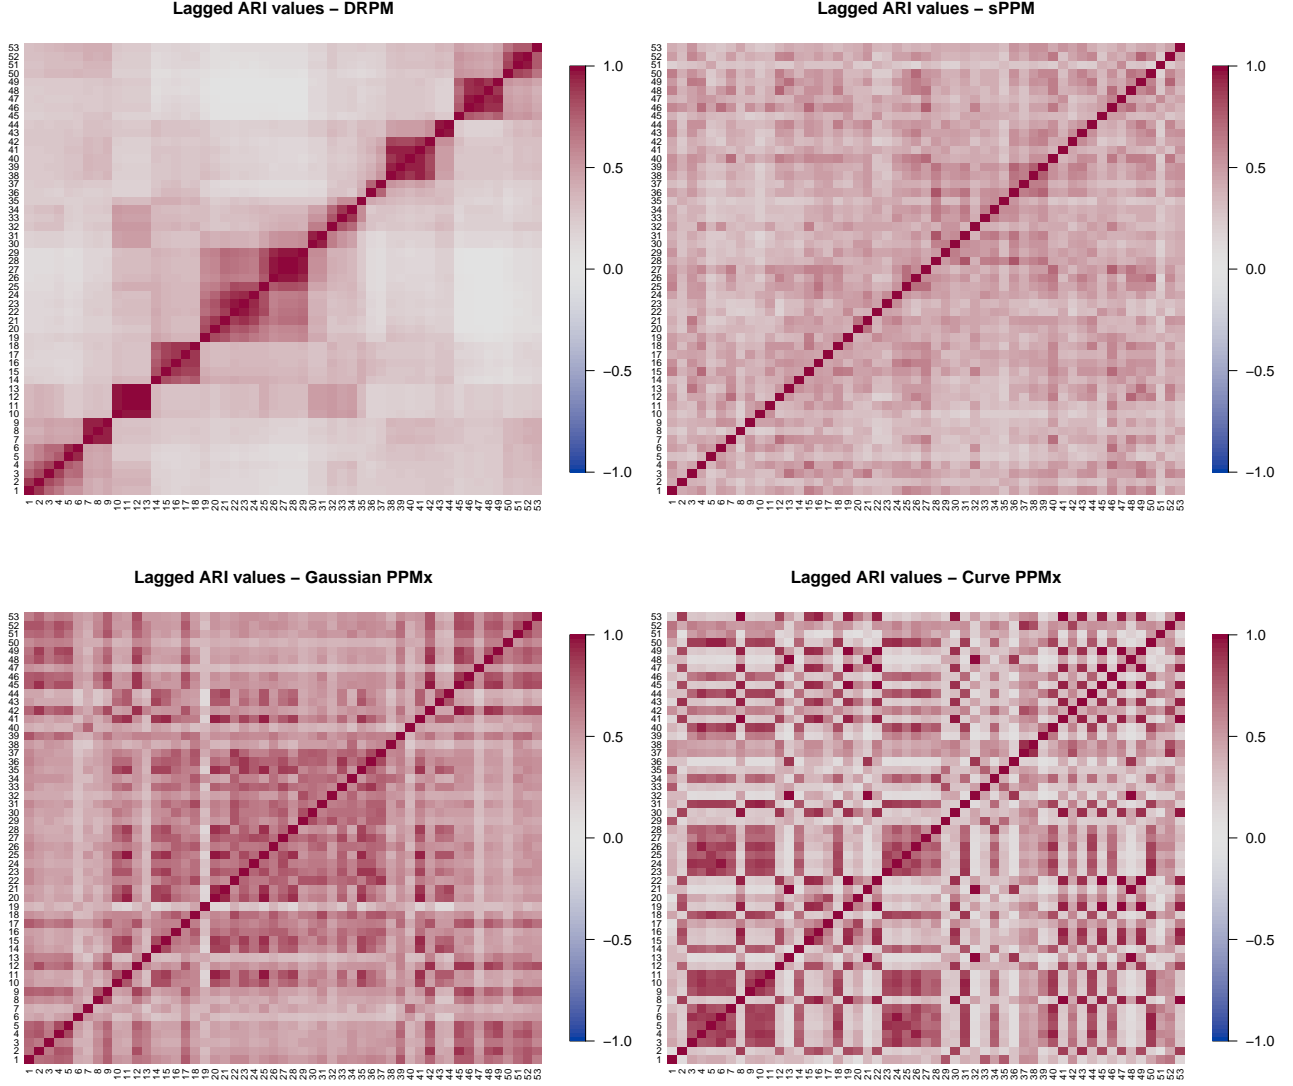
Figure 5: Lagged ARI values of the four models.

when the two units are clustered together in both ρ_1 and ρ_2 allocations) and b the number of pairs allocated in different clusters (i.e. when the two units do not belong to the same cluster in both ρ_1 and ρ_2 allocations).

This metric allows for example to compare a proposal clustering with the real one, if available, to see how good is the matching; but in our case, where there was no correct answer, we used it to study the time evolution of the clusterings and to check the agreement level among the different models.

Regarding the time evolution we studied the Lagged ARI values, meaning that for each model we computed $\mathrm{ARI}(\rho_t, \rho_{t+k})$ for $t \in \{1, \dots, 53\}$ and for all the valid values of k. In this way we obtained an information about the relation among the clusters throughout the year. As depicted in Figure 5, we can see how the DRPM exhibits a gentle evolution of the clusters, in a sort of time persistency, where almost every ρ_t tends to be similar to the subsequents clusterings, affecting the ρ_{t+k} for k around 2 or 3 time instants, while losing connections with the ones of further away in time. The sPPM and Gaussian PPMx instead do not seem to show any particular pattern, expect for the latter which displays a general dominance of high values, denoting maybe a more rigid and recurrent structure. Also Curve PPMx shows an interesting but quite erratic pattern, with an appearance of some spots where clusters tend to remain similar to each other, like it happened in the DRPM case, but now with less regularity.

The other aspect that the ARI metric allowed to study was the agreement among the clusterings generated by the models, to see if a sort of general consensus and division appeared clearly in all models. To study this we computed $\mathrm{ARI}(\rho_t^{M_1}, \rho_t^{M_2})$ for $t \in \{1, \dots, 53\}$ and for each pair of models M_1 and M_2 in the four that we fitted. We can see that ...

Figure 6: Plot of the ARI values for each pairwise comparison of the models, for all the weeks.

5. Analysis of the results

6. Conclusions

7. Further developements

Several avenues for a further development remain, presenting opportunities to enhance the depth and precision of our analysis:

- *Utilize Previous Year Data for Model Priors:* consider incorporating data from the preceding year to establish priors for the models. While our current dataset facilitated model convergence, integrating historical data could offer additional insights and refine the robustness of our findings.
- *Distinguish Between Weekends and Weekdays:* exploring the impact of human activities on PM10 levels by stratifying the analysis between weekends and weekdays. This differentiation may uncover patterns associated with specific human-related factors, contributing to a more nuanced understanding of particulate matter dynamics.
- *Ensemble Modeling:* exploring the potential benefits of ensemble modeling by combining the outputs of different models. This approach can enhance the overall accuracy and reliability of our clustering analysis. By leveraging the strengths of individual models, we can obtain a more comprehensive and robust estimation of the identified clusters.

These proposed extensions aim to further refine our methodology, enrich the interpretability of results, and provide a more comprehensive understanding of the intricate factors influencing PM10 levels in the Lombardy region.

References

[Com17]   European Commission. Infringement actions for excessive levels of PM10 in Italy, 2017. `https://ec.europa.eu/commission/presscorner/detail/ET/IP_17_1046`.

[CTAa]   CTAN. BiBTeX documentation.

[CTAb]   CTAN. pgf – create PostScript and PDF graphics in TEX.

[FRFM⁺23] A. Fassò, J. Rodeschini, A. Fusta Moro, Q. Shaboviq, P. Maranzano, M. Cameletti, F. Finazzi, N. Golini, R. Ignaccolo, and P. Otto. AgrImOnIA: Open Access dataset correlating livestock and air quality in the Lombardy region, Italy (3.0.0), 2023.

[HA85]   Lawrence J. Hubert and Phipps Arabie. Comparing partitions. *Journal of Classification*, 2:193–218, 1985.

[Knu74]   Donald E. Knuth. Computer programming as an art. *Commun. ACM*, pages 667–673, 1974.

[Knu92]   Donald E. Knuth. Two notes on notation. *Amer. Math. Monthly*, 99:403–422, 1992.

[Kot15]   Stefan Kottwitz. *LaTeX Cookbook*. Packt Publishing Ltd, 2015.

[Lam94]   Leslie Lamport. *LaTeX: A Document Preparation System*. Pearson Education India, 1994.

[OPHS95] Tobias Oetiker, Hubert Partl, Irene Hyna, and Elisabeth Schlegl. The not so short introduction to latex2ε. *Electronic document available at http://www. tex. ac. uk/tex-archive/info/lshort*, 1995.

[PQD22]   Garrit L. Page, Fernando A. Quintana, and David B. Dahl. Dependent modeling of temporal sequences of random partitions. *Journal of Computational and Graphical Statistics*, 31(2):614–627, 2022.

A.   Appendix A

A pivotal component of a robust statistical analysis lies in the effective interpretation of results. To address this crucial aspect, we meticulously constructed a library of auxiliary functions, empowering us to visually scrutinize various facets of our research.

Given the inherent temporal and spatial dimensions of our dataset, we opted for a dynamic approach, creating videos instead of static images to seamlessly navigate the temporal component.

For the visualisation of spatial variables, we devised two principal tools: a grid map and an expanding circles plot.

1. **Grid Map:** - This tool harnesses a distinct dataset featuring measurements across the entire region, organized on a grid of evenly spaced points. It offers a panoramic overview of key variables, such as Altitude and Weather measurements, providing a comprehensive understanding of spatial patterns. 2. **Expanding Circles Plot:** - Focused on station-level measurements, this tool illustrates the magnitude of variables by employing radius and color intensity of circles centered around each station. This approach grants us insights into localized patterns, enhancing our comprehension of variable distribution across the region.

To enhance the clarity of our cluster representations, we devised a function that establishes connections between stations within the same cluster. These connections are formed by solving a minimum spanning tree, a strategic approach chosen to yield a more organized and visually coherent representation of the clusters.

B.   MCMC diagnostics DONE

Here we present the plots that we used to check the convergence of the MCMC values generated by the models fit functions. They are, for summary purposes, just on two of the most relevant parameters of each model, and on weeks running from 1 to 50 with jumps of length 7.

We used trace plots to ensure that the chosen values for the burn-in were high enough to remove any unstable behaviour in the iterates. However, even after really high burn-in periods (e.g. 60000 in DRPM), occasionally there were still some oscillations, but by looking at the y axis we can see that there is no significant variation after all. Also sometimes there are divergent iterations, especially in the σ^2 trace plots. We think that these

small issues were due to the complexity of the models, spanning on lots of subjects (the 105 stations) and several time instants (the 53 weeks), together with implementing a deep hierarchical structure.

We also looked at ACF plots to tune the thinning parameter, by seeing the trend of the auto correlation on subsequent iterates.
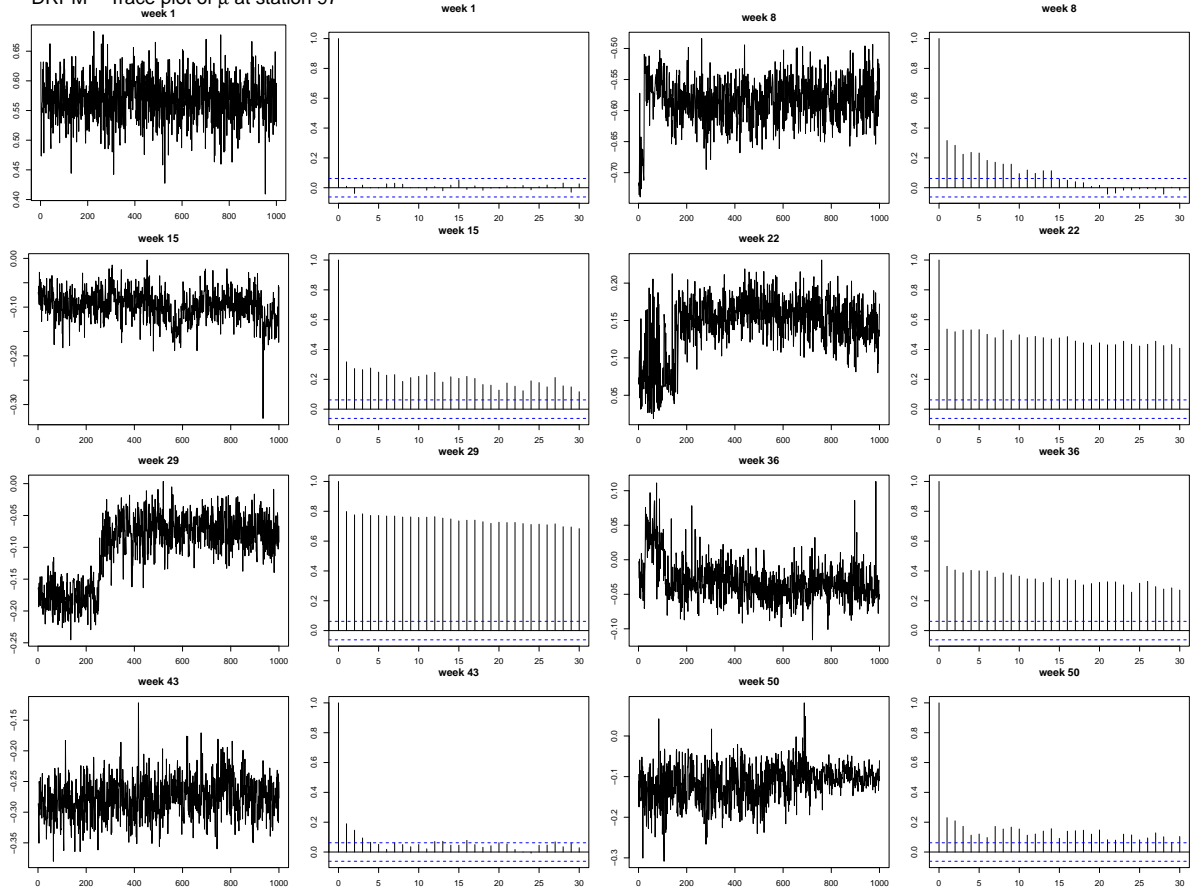
Figure 7: Trace and ACF plots of parameters σ^2 and μ of the DRPM model.

13

sPPM – Trace plot of σ² at station 35

sPPM – Trace plot of μ at station 35

Figure 8: Trace and ACF plots of parameters σ^2 and μ of the sPPM model.

14
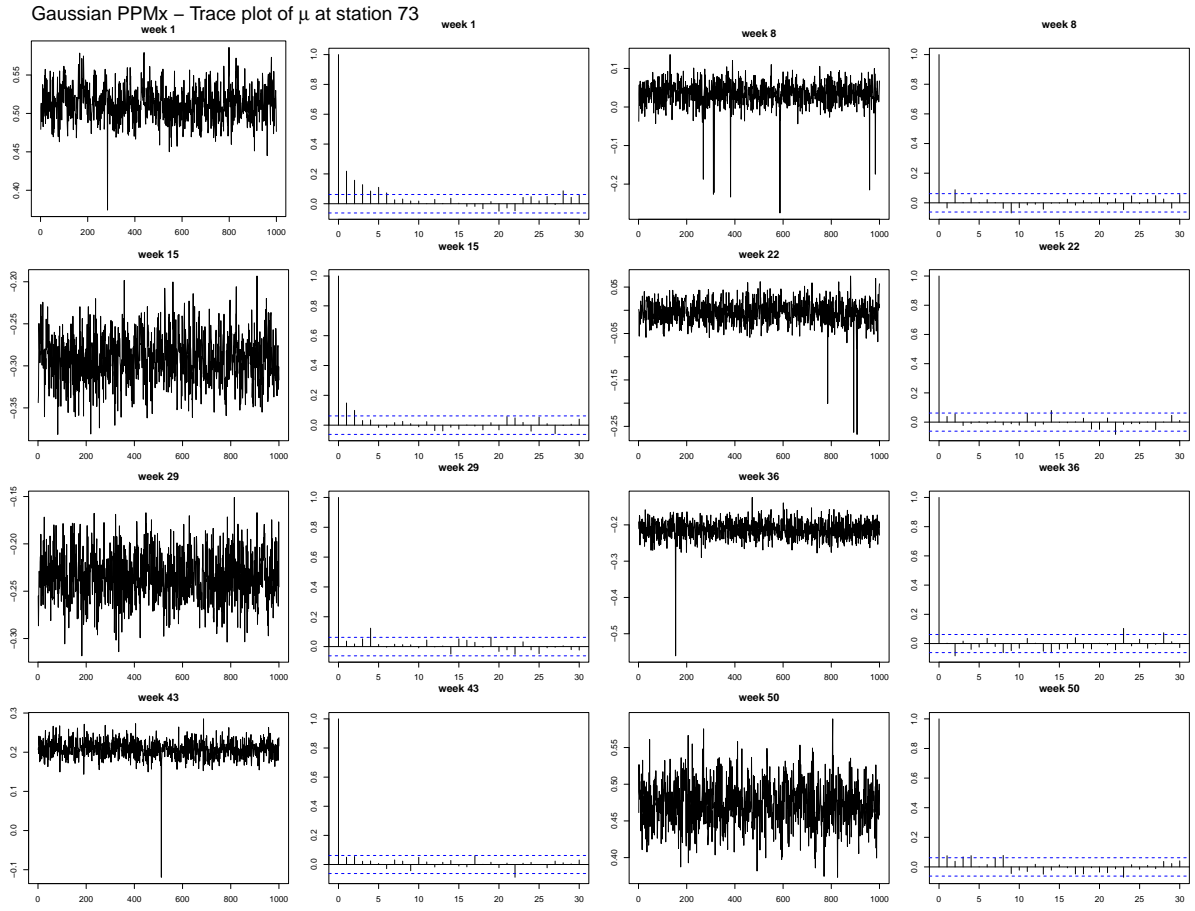
Figure 9: Trace and ACF plots of parameters σ^2 and μ of the Gaussian PPMx model.
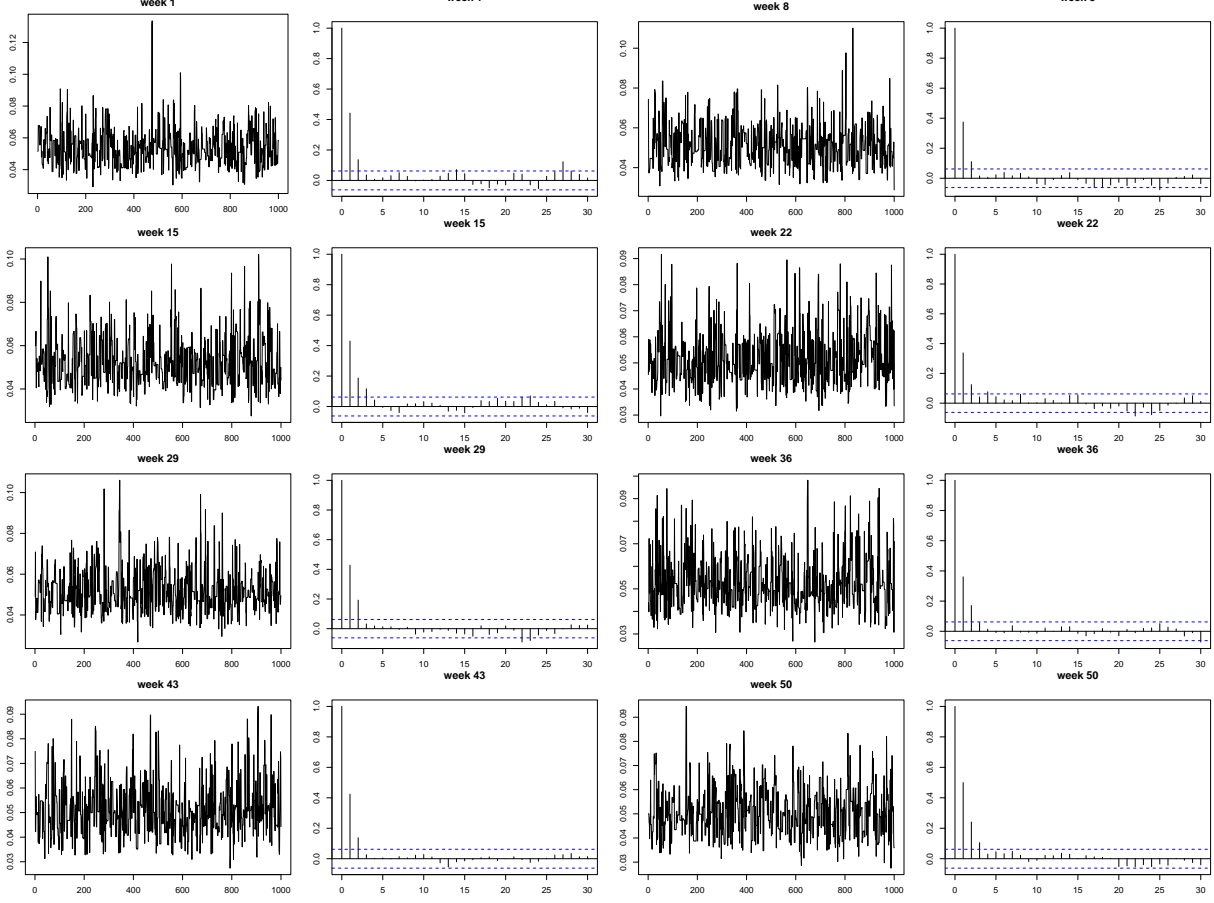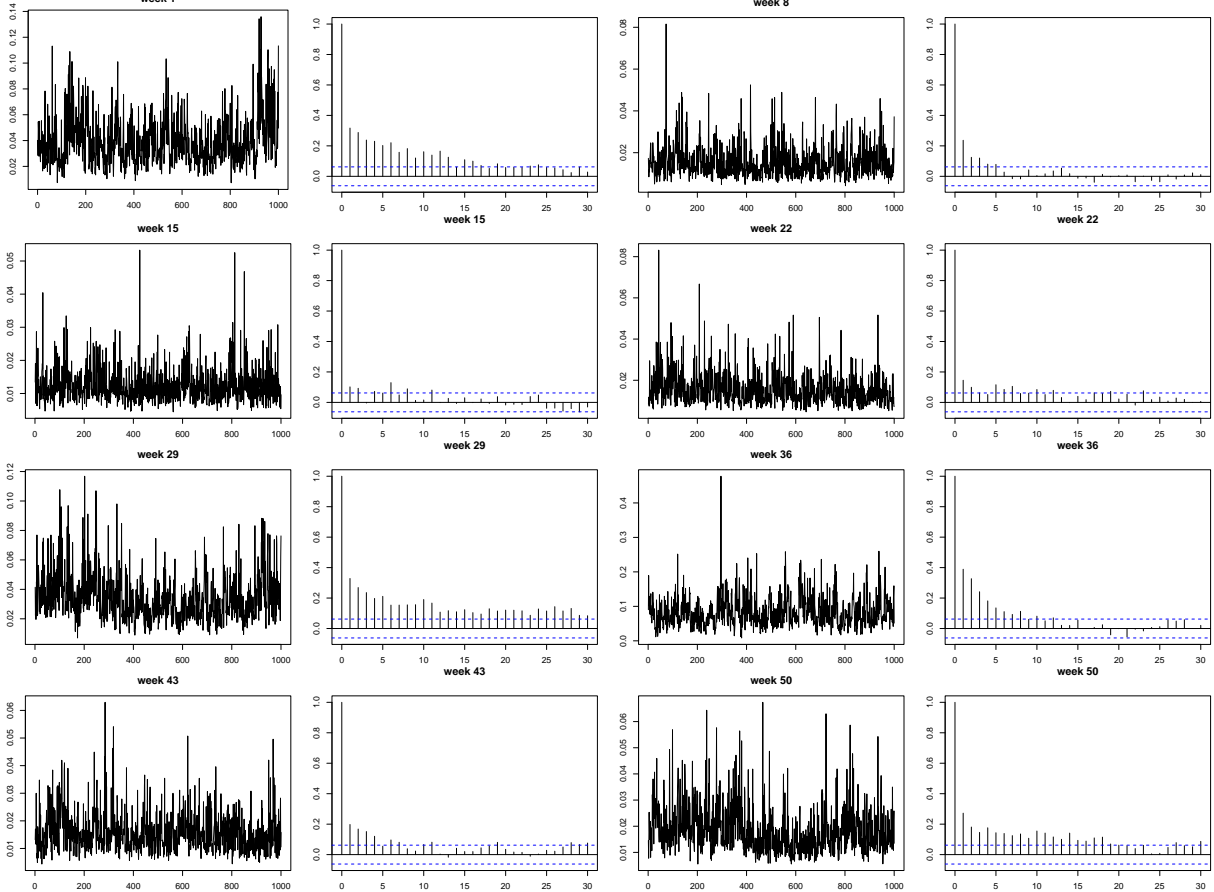
Figure 10: Trace and ACF plots of parameters σ^2 and τ^2 of the Curve PPMx model.

16

===========================

POLIMI TEMPLATE EXAMPLE

C.   Introduction

This document is intended to be both an example of the Polimi LaTeX template for Master Theses in article format, as well as a short introduction to its use. It is not intended to be a general introduction to LaTeX itself, and the reader is assumed to be familiar with the basics of creating and compiling LaTeX documents (see [OPHS95, Kot15]).

The cover page of the thesis in article format must contain all the relevant information: title of the thesis, name of the Study Programme, name(s) of the author(s), student ID number, name of the supervisor, name(s) of the co-supervisor(s) (if any), academic year.

Be sure to select a title that is meaningful. It should contain important keywords to be identified by indexer. Keep the title as concise as possible and comprehensible even to people who are not experts in your field. The title has to be chosen at the end of your work so that it accurately captures the main subject of the manuscript. It is convenient to break the article format of your thesis (in article format) into sections and subsections. If necessary, subsubsections, paragraphs and subparagraphs can be used. A new section is created by the command

`\section{Title of the section}`

The numbering can be turned off by using `\section*{}`. A new subsection is created by the command

`\subsection{Title of the subsection}`

and, similarly, the numbering can be turned off by adding an asterisk as follows

`\subsection*{}`

It is recommended to give a label to each section by using the command

`\label{sec:section_name}%`

where the argument is just a text string that you'll use to reference that part as follows: *Section C contains* INTRODUCTION . . . .

D.   Equations

This section gives some examples of writing mathematical equations in your thesis.
Maxwell's equations read:

$$\begin{cases} \nabla \cdot \boldsymbol{D} = \rho, & \text{(2a)} \\ \nabla \times \boldsymbol{E} + \dfrac{\partial \boldsymbol{B}}{\partial t} = \boldsymbol{0}, & \text{(2b)} \\ \nabla \cdot \boldsymbol{B} = 0, & \text{(2c)} \\ \nabla \times \boldsymbol{H} - \dfrac{\partial \boldsymbol{D}}{\partial t} = \boldsymbol{J}. & \text{(2d)} \end{cases}$$

Equation (2) is automatically labeled by `cleveref`, as well as Equation (2a) and Equation (2c). Thanks to the `cleveref` package, there is no need to use `\eqref`. Equations have to be numbered only if they are referenced in the text.

Equations (3), (4), (5), and (6) show again Maxwell's equations without brace:

$$\nabla \cdot \boldsymbol{D} = \rho, \tag{3}$$

$$\nabla \times \boldsymbol{E} + \frac{\partial \boldsymbol{B}}{\partial t} = \boldsymbol{0}, \tag{4}$$

$$\nabla \cdot \boldsymbol{B} = 0, \tag{5}$$

$$\nabla \times \boldsymbol{H} - \frac{\partial \boldsymbol{D}}{\partial t} = \boldsymbol{J}. \tag{6}$$

Equation (7) is the same as before, but with just one label:

$$\begin{cases} \nabla \cdot \boldsymbol{D} = \rho, \\ \nabla \times \boldsymbol{E} + \dfrac{\partial \boldsymbol{B}}{\partial t} = \boldsymbol{0}, \\ \nabla \cdot \boldsymbol{B} = 0, \\ \nabla \times \boldsymbol{H} - \dfrac{\partial \boldsymbol{D}}{\partial t} = \boldsymbol{J}. \end{cases} \tag{7}$$

E. Figures, Tables and Algorithms

Figures, Tables and Algorithms have to contain a Caption that describes their content, and have to be properly referred in the text.

E.1. Figures

For including pictures in your text you can use `TikZ` for high-quality hand-made figures [CTAb], or just include them with the command

`\includegraphics[options]{filename.xxx}`

Here xxx is the correct format, e.g. `.png`, `.jpg`, `.eps`, ...

Figure 11: Caption of the Figure.

Thanks to the `\subfloat` command, a single figure, such as Figure 11, can contain multiple sub-figures with their own caption and label, e.g. Figure 12a and Figure 12b.

(a) One PoliMi logo. (b) Another one PoliMi logo.

Figure 12: Caption of the Figure.

E.2. Tables

Within the environments `table` and `tabular` you can create very fancy tables as the one shown in Table 3.

Example of Table (optional)

	column1	column2	column3
row1	1	2	3
row2	α	β	γ
row3	alpha	beta	gamma

Table 3: Caption of the Table.

You can also consider to highlight selected columns or rows in order to make tables more readable. Moreover, with the use of `table*` and the option `bp` it is possible to align them at the bottom of the page. One example is presented in Table 4.

E.3. Algorithms

Pseudo-algorithms can be written in LATEX with the `algorithm` and `algorithmic` packages. An example is shown in Algorithm 1.

Algorithm 1 Name of the Algorithm

1: Initial instructions
2: **for** $for - condition$ **do**
3:     Some instructions
4:     **if** $if - condition$ **then**
5:         Some other instructions
6:     **end if**
7: **end for**
8: **while** $while - condition$ **do**
9:     Some further instructions
10: **end while**
11: Final instructions

F. Some further useful suggestions

Theorems have to be formatted as follows:

Theorem F.1. *Write here your theorem.*

Proof. If useful you can report here the proof.

Propositions have to be formatted as follows:

Proposition F.1. *Write here your proposition.*

How to insert itemized lists:
- first item;

	column1	column2	column3	column4	column5	column6
row1	1	2	3	4	5	6
row2	a	b	c	d	e	f
row3	α	β	γ	δ	ϕ	ω
row4	alpha	beta	gamma	delta	phi	omega

Table 4: Highlighting the columns

- second item.

How to write numbered lists:

1. first item;
2. second item.

G. Use of copyrighted material

Each student is responsible for obtaining copyright permissions, if necessary, to include published material in the thesis. This applies typically to third-party material published by someone else.

H. Plagiarism

You have to be sure to respect the rules on Copyright and avoid an involuntary plagiarism. It is allowed to take other persons' ideas only if the author and his original work are clearly mentioned. As stated in the Code of Ethics and Conduct, Politecnico di Milano *promotes the integrity of research, condemns manipulation and the infringement of intellectual property*, and gives opportunity to all those who carry out research activities to have an adequate training on ethical conduct and integrity while doing research. To be sure to respect the copyright rules, read the guides on Copyright legislation and citation styles available at:

`https://www.biblio.polimi.it/en/tools/courses-and-tutorials`

You can also attend the courses which are periodically organized on "Bibliographic citations and bibliography management".

I. Conclusions

A final section containing the main conclusions of your research/study and possible future developments of your work have to be inserted in the section "Conclusions".

J. Bibliography and citations

Your thesis must contain a suitable Bibliography which lists all the sources consulted on developing the work. The list of references is placed at the end of the manuscript after the chapter containing the conclusions. It is suggested to use the BibTeX package and save the bibliographic references in the file `bibliography.bib`. This is indeed a database containing all the information about the references. To cite in your manuscript, use the `\cite{}` command as follows:
Here is how you cite bibliography entries: [Knu74], or multiple ones at once: [Knu92, Lam94].
The bibliography and list of references are generated automatically by running BibTeX [CTAa].